

# Jyothi Vardhana Rao Metta

Tampa, FL

+1-813-797-9106 | vardhanrao9849@gmail.com | LinkedIn | GitHub | LeetCode

## Summary

Software Engineer— AWS Glue | Redshift | Snowflake who architects Petabyte-scale pipelines that trim refresh latency 70 % and cloud spend 22 %. 4 years delivering real-time analytics to 500 + users; AWS DEA & DVA certified; Power BI / Tableau storyteller.

## TECHNICAL SKILLS

- **Data Processing:** Python(Pandas, NumPy, scikit-learn), SQL, PySpark, R, Scala, DBT,Matlab, Excel, Hadoop
- **Cloud & DevOps:** AWS(S3, Lambda, Glue, EMR, EC2, IAM, EMR, Firehose, Kinesis), Azure Data Factory, Docker, Snowflake, Kubernetes, BigQuery
- **Databases:** PostgreSQL, MongoDB, Oracle Database, Redshift, DynamoDB, Redis
- **Orchestration & IaC:** Airflow, Step Functions, Kafka, Terraform, GitHub Actions, Flink, Jenkins
- **BI / Analytics:** PowerBI, Tableau, QuickSight, Looker, Qlik

## EXPERIENCE

Infosys Pvt Ltd. (Client-NextEra Energy Inc.)

Dec 2021 - Jul 2023

Software Engineer

India

- Deployed 7+ real-time, **cost-optimized pipelines** with **AWS Lambda**, Amazon **MSK**, and **API Gateway** to replace on-prem consumer jobs, increasing data-processing capacity and documenting the architecture in **Confluence** for seamless cross-team adoption.
- Engineered **PySpark ETL** in **AWS Glue** to transform high-volume **Parquet** data from **S3** into **Redshift**, boosting pipeline throughput and query performance by **20 %** for downstream **analytics**.
- Architected 25 event-driven **DAGs** in **Apache Airflow** on **ECS Fargate**—triggered via **Lambda**, Step Functions, and **S3** metadata sensors—to automate multi-stage refreshes and slash end-to-end data-latency 70 %.
- Locked down 12 enterprise **dashboards** by implementing dynamic **Row-Level Security (RLS)** with **DAX** and role-based access for 600 users across **7 permission groups**, cutting unauthorized-access.
- Migrated **42 TB** of **Snowflake** data to Amazon **S3** with **AWS DMS**, auto-cataloged via **Glue**, verified in **Athena**, and trimmed storage costs 28 %.
- Spearheaded cross-functional initiatives for **NextEra Energy**, orchestrating daily **Scrum** calls with 5-person squad and **Jira**-based sprint planning to institutionalize **Agile** adoption and boost workflow efficiency 15 %.

Digital IT Hub

Jan 2020 - Nov 2021

Data Engineer

India

- Authored 30+ **CloudFormation** stacks (YAML) linked to **GitHub Actions CI/CD**, enabling version-controlled **dev/QA/prod** provisioning that spins environments up 5× faster and ends config drift.
- Developed modular **DBT SQL** models on **MySQL** with ref()-driven **DAG** dependencies and built-in tests to enable scalable **ETL** transformations for analytics, boosted performance by 35 % through incremental materializations and targeted **SQL** query optimizations.
- Built **Python** scripts in **Azure Databricks** to ingest and transform Parquet, CSV, and JSON datasets, optimizing **Spark** configurations and dynamic cluster scaling to boost processing capacity by 200 % and cut runtime by 60 %.

Bharat Dynamics Limited (Government of India Enterprise)

Feb 2019 - Jul 2019

Intern

India

- Collaborated with senior scientists to preprocess and standardize 50 GB of complex geo-location data using **Python** scripts and **pandas** for accurate downstream analysis.
- Optimized **SQL** and **PL/SQL** procedures on **Oracle** databases, reducing execution time by 45 %, and managed installations and configurations of OBIA 10.2 and OBIEE 12c for enterprise reporting.
- Monitored 15+ **Kafka** consumer jobs, tracking lag/throughput to ensure 99.9 % uptime.

## EDUCATION

University of South Florida, Muma College of Business — Tampa, FL, USA

Aug 2023 - May 2025

Master of Science, Business Analytics and Information Systems

- **GPA:** 3.8/4.0

## PROJECTS

Analysis of Seasonal Hash Rate Changes and Crypto Price Impact

Feb 2025 - Mar 2025

- Conducted Time Series Analysis on historical crypto price data using **AWS Glue**, **Athena**, **Lambda**, **Kafka** and **S3**, and built a **Tableau** dashboard highlighting trends and potential correlations with temperature-spotlighting Binance Coin's 133 % growth.

Real-time Custom Weather API

Mar 2025 - Apr 2025

- Built a Real-Time weather data pipeline using OpenWeather, **Kinesis**, and **ECS Fargate**, storing processed data in **DynamoDB** and exposing it via secure, versioned endpoints through **API Gateway** for querying U.S. state-level weather conditions.
- Created a comprehensive **Swagger** specification document to standardize and document the **API** for developer usability

## CERTIFICATIONS

- AWS Certified Developer - Associate (Aug 2024 - Aug 2027)
- AWS Certified Data Engineer - Associate (Jan 2025 - Jan 2028)